

Homework 3

Naga Kartheek Peddisetty, 50538422

10/15/2023

Question 1) Consider the Diabetes dataset (posted with assignment). Assume the population prior probabilities are estimated using the relative frequencies of the classes in the data.

```
library(MASS)
library(ggplot2)
setwd("D:/Buffalo/files")
load("D:/Buffalo/files/Diabetes.RData")
dim(Diabetes) # 145 * 6
```

```
## [1] 145 6
```

```
str(Diabetes)
```

```
## 'data.frame': 145 obs. of 6 variables:
## $ relwt : num 0.81 0.95 0.94 1.04 1 0.76 0.91 1.1 0.99 0.78 ...
## $ glufast: int 80 97 105 90 90 86 100 85 97 97 ...
## $ glutest: int 356 289 319 356 323 381 350 301 379 296 ...
## $ instest: int 124 117 143 199 240 157 221 186 142 131 ...
## $ sspg : int 55 76 105 108 143 165 119 105 98 94 ...
## $ group : Factor w/ 3 levels "Normal","Chemical_Diabetic",...: 1 1 1 1 1 1 1 1 1 1 ...
```

```
head(Diabetes)
```

```
## relwt glufast glutest instest sspg group
## 1 0.81 80 356 124 55 Normal
## 2 0.95 97 289 117 76 Normal
## 3 0.94 105 319 143 105 Normal
## 4 1.04 90 356 199 108 Normal
## 5 1.00 90 323 240 143 Normal
## 6 0.76 86 381 157 165 Normal
```

```
table(Diabetes$group)
```

```
##
## Normal Chemical_Diabetic Overt_Diabetic
## 76 36 33
```

```
diabetes <- Diabetes
```

```
### Converting int to numeric data
```

```
diabetes[1:5] <- data.frame(sapply(diabetes[1:5], as.numeric))  
str(diabetes)
```

```
## 'data.frame': 145 obs. of 6 variables:  
## $ relwt : num 0.81 0.95 0.94 1.04 1 0.76 0.91 1.1 0.99 0.78 ...  
## $ glufast: num 80 97 105 90 90 86 100 85 97 97 ...  
## $ glutest: num 356 289 319 356 323 381 350 301 379 296 ...  
## $ instest: num 124 117 143 199 240 157 221 186 142 131 ...  
## $ sspg : num 55 76 105 108 143 165 119 105 98 94 ...  
## $ group : Factor w/ 3 levels "Normal","Chemical_Diabetic",...: 1 1 1 1 1 1 1 1 1 1 ...
```


```
cor(diabetes[1:5])
```

```
##           relwt      glufast      glutest      instest      sspg  
## relwt      1.000000000 -0.008813193  0.0239843  0.222237813  0.384319804  
## glufast -0.008813193  1.000000000  0.9646281 -0.396234858  0.715480192  
## glutest  0.023984304  0.964628091  1.0000000 -0.337020435  0.770942459  
## instest  0.222237813 -0.396234858 -0.3370204  1.000000000  0.007914263  
## sspg      0.384319804  0.715480192  0.7709425  0.007914263  1.000000000
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
corrplot(cor(diabetes[1:5]), tl.col = "blue", tl.srt= 45, tl.cex=1, cl.cex=1)
```



a) Produce pairwise scatterplots for all five variables, with different symbols or colors representing the three different classes. Do you see any evidence that the classes may have difference covariance matrices? That they may not be multivariate normal?

```
library(psych)
```

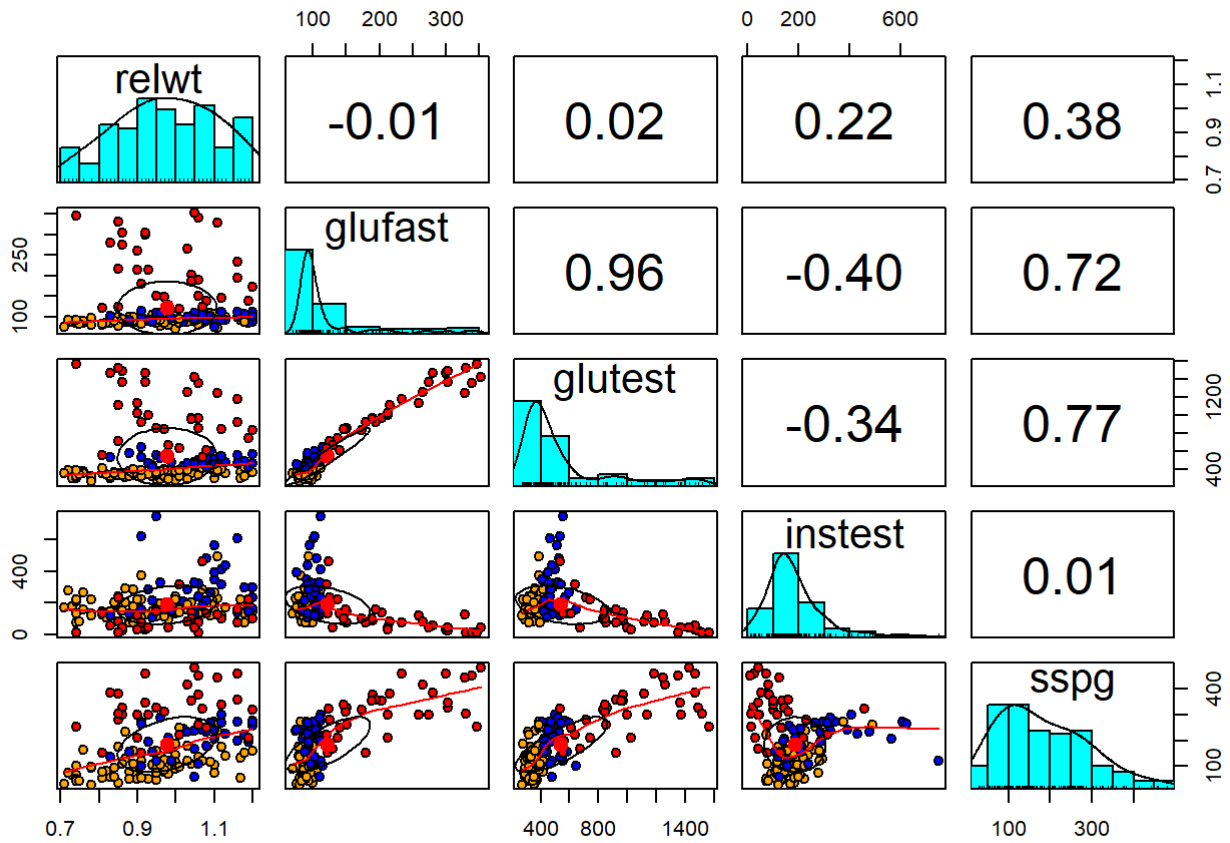
```
##  
## Attaching package: 'psych'
```

```
## The following objects are masked from 'package:ggplot2':  
##  
## %+%, alpha
```

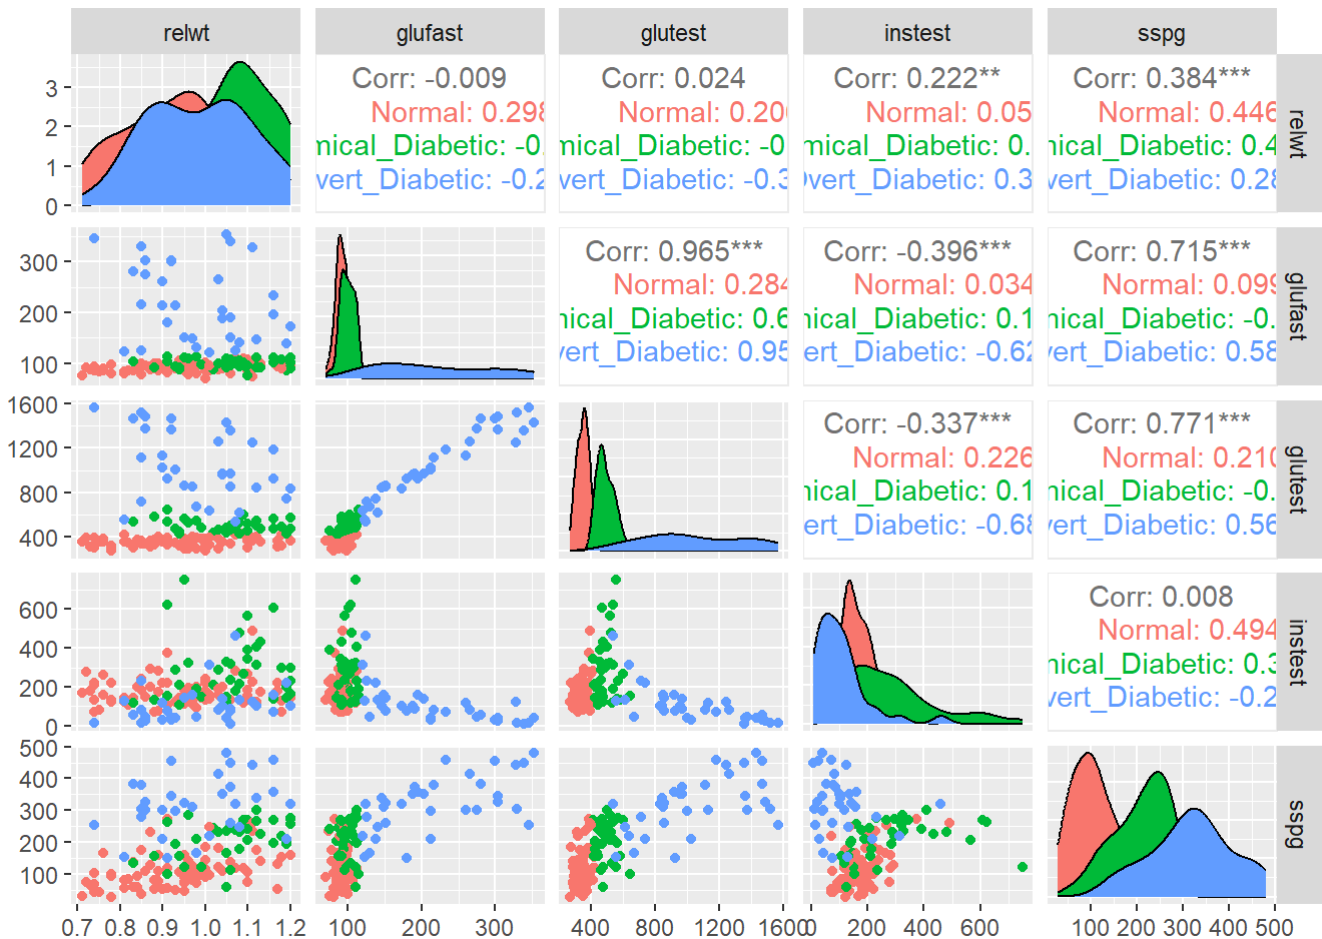
```
library(GGally)
```

```
## Registered S3 method overwritten by 'GGally':  
## method from  
## +.gg ggplot2
```

```
pairs.panels(diabetes[,1:5],  
             bg = c("orange", "blue", "red")[diabetes$group],  
             pch = 21)
```



```
ggpairs(diabetes[, 1:5], aes(color = diabetes$group))
```



From the above plot it is clear that classes have different covariance matrices. Because, the plot shows the different sizes of clusters or shapes for different classes.

Classes are also multivariate normal(normal distribution can be observed in the above plots).

b) Apply linear discriminant analysis (LDA) and quadratic discriminant analysis (QDA). How does the performance of QDA compare to that of LDA in this case?

```
### Creating a test and training data
```

classes are not multivariate normal; 2 or more variables are to be considered to check for the condition of multivariate normality (-1 mark)

```
set.seed(123)
```

```
indis <- sample(1:nrow(diabetes),size = round(0.7 * nrow(diabetes)))
```

```
train_data <- diabetes[indis, ]
```

```
test_data <- diabetes[-indis, ]
```

```
X_train <- train_data[, -6]
```

```
Y_train <- train_data[, 6]
```

```
X_test <- test_data[, -6]
```

```
Y_test <- test_data[,6]
```

LDA

```
lda.fit <- lda(group ~., data = train_data)
```

```
lda.fit
```

```
## Call:
```

```
## lda(group ~ ., data = train_data)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##           Normal Chemical_Diabetic   Overt_Diabetic
```

```
##           0.5196078         0.2058824         0.2745098
```

```
##
```

```
## Group means:
```

```
##           relwt  glufast  glutest  instest  sspg
```

```
## Normal           0.9364151  91.0000  350.8679 184.3585 124.0566
```

```
## Chemical_Diabetic 1.0823810 100.8095  503.1905 323.4286 219.5714
```

```
## Overt_Diabetic   0.9775000 214.3571 1029.3214 103.4286 315.6786
```

```
##
```

```
## Coefficients of linear discriminants:
```

```
##           LD1          LD2
```

```
## relwt    1.8137899322 -4.824493227
```

```
## glufast  -0.0301791937  0.026377497
```

```
## glutest  0.0115649319 -0.005711114
```

```
## instest  -0.0004145498 -0.006564432
```

```
## sspg      0.0034478417  0.002125798
```


```
##
```

```
## Proportion of trace:
```

```
##    LD1    LD2
```

```
## 0.8603 0.1397
```

```
plot(lda.fit)
```



```
# predictions for the test and training.
```

```
test_pred_lda <- predict(lda.fit, newdata = test_data)
names(test_pred_lda)
```

```
## [1] "class"      "posterior" "x"
```

```
test_pred_lda$class
```

```
## [1] Normal      Normal      Normal      Normal
## [5] Normal      Normal      Normal      Normal
## [9] Normal      Normal      Normal      Normal
## [13] Normal     Normal      Normal      Normal
## [17] Normal     Normal      Normal      Normal
## [21] Chemical_Diabetic Chemical_Diabetic Normal      Normal
## [25] Normal     Chemical_Diabetic Normal      Normal
## [29] Chemical_Diabetic Chemical_Diabetic Normal      Chemical_Diabetic
## [33] Chemical_Diabetic Chemical_Diabetic Chemical_Diabetic Normal
## [37] Chemical_Diabetic Normal      Overt_Diabetic Overt_Diabetic
## [41] Overt_Diabetic Overt_Diabetic Chemical_Diabetic
## Levels: Normal Chemical_Diabetic Overt_Diabetic
```

```
train_pred_lda <- predict(lda.fit, newdata = train_data)
train_pred_lda$class
```

```
## [1] Normal Normal Overt_Diabetic Normal
## [5] Overt_Diabetic Overt_Diabetic Chemical_Diabetic Chemical_Diabetic
## [9] Overt_Diabetic Chemical_Diabetic Chemical_Diabetic Normal
## [13] Chemical_Diabetic Normal Normal Normal
## [17] Overt_Diabetic Normal Overt_Diabetic Normal
## [21] Normal Normal Chemical_Diabetic Chemical_Diabetic
## [25] Overt_Diabetic Normal Normal Normal
## [29] Normal Normal Normal Normal
## [33] Overt_Diabetic Overt_Diabetic Overt_Diabetic Normal
## [37] Normal Chemical_Diabetic Normal Chemical_Diabetic
## [41] Normal Normal Overt_Diabetic Chemical_Diabetic
## [45] Normal Chemical_Diabetic Chemical_Diabetic Chemical_Diabetic
## [49] Normal Overt_Diabetic Normal Normal
## [53] Overt_Diabetic Normal Overt_Diabetic Overt_Diabetic
## [57] Normal Normal Chemical_Diabetic Chemical_Diabetic
## [61] Normal Normal Overt_Diabetic Overt_Diabetic
## [65] Chemical_Diabetic Normal Chemical_Diabetic Normal
## [69] Overt_Diabetic Normal Normal Chemical_Diabetic
## [73] Normal Normal Normal Normal
## [77] Normal Normal Normal Normal
## [81] Overt_Diabetic Normal Overt_Diabetic Overt_Diabetic
## [85] Normal Overt_Diabetic Normal Normal
## [89] Overt_Diabetic Chemical_Diabetic Normal Normal
## [93] Normal Normal Chemical_Diabetic Normal
## [97] Normal Chemical_Diabetic Overt_Diabetic Normal
## [101] Chemical_Diabetic Normal
## Levels: Normal Chemical_Diabetic Overt_Diabetic
```

```
# compute the error rates
```

```
train_error_lda <- (1/length(train_pred_lda$class))*length(which(Y_train != train_pred_lda$class))
test_error_lda <- (1/length(test_pred_lda$class))*length(which(Y_test != test_pred_lda$class))

train_error_lda
```

```
## [1] 0.09803922
```

```
test_error_lda
```

```
## [1] 0.1395349
```

QDA

```
qda.fit <- qda(group ~., data = train_data)
qda.fit
```

```
## Call:
## qda(group ~ ., data = train_data)
##
## Prior probabilities of groups:
##           Normal Chemical_Diabetic   Overt_Diabetic
##           0.5196078       0.2058824       0.2745098
##
## Group means:
##           relwt  glufast  glutest  instest    sspg
## Normal           0.9364151  91.0000  350.8679 184.3585 124.0566
## Chemical_Diabetic 1.0823810 100.8095  503.1905 323.4286 219.5714
## Overt_Diabetic   0.9775000 214.3571 1029.3214 103.4286 315.6786
```

```
# predictions for the test and training.
```

```
test_pred_qda <- predict(qda.fit, newdata = test_data)
class(test_pred_qda)
```

```
## [1] "list"
```

```
names(test_pred_qda)
```

```
## [1] "class"      "posterior"
```

```
train_pred_qda <- predict(qda.fit, newdata = train_data)
train_pred_qda$class
```



```
## [1] Normal Normal Overt_Diabetic Normal
## [5] Overt_Diabetic Overt_Diabetic Chemical_Diabetic Chemical_Diabetic
## [9] Overt_Diabetic Chemical_Diabetic Chemical_Diabetic Normal
## [13] Normal Normal Normal Normal
## [17] Overt_Diabetic Chemical_Diabetic Overt_Diabetic Normal
## [21] Normal Normal Chemical_Diabetic Chemical_Diabetic
## [25] Overt_Diabetic Normal Normal Normal
## [29] Normal Normal Normal Normal
## [33] Overt_Diabetic Overt_Diabetic Overt_Diabetic Chemical_Diabetic
## [37] Normal Chemical_Diabetic Normal Chemical_Diabetic
## [41] Normal Overt_Diabetic Overt_Diabetic Chemical_Diabetic
## [45] Normal Chemical_Diabetic Chemical_Diabetic Overt_Diabetic
## [49] Normal Overt_Diabetic Normal Normal
## [53] Overt_Diabetic Normal Overt_Diabetic Overt_Diabetic
## [57] Normal Normal Chemical_Diabetic Chemical_Diabetic
## [61] Normal Normal Overt_Diabetic Overt_Diabetic
## [65] Overt_Diabetic Normal Chemical_Diabetic Chemical_Diabetic
## [69] Overt_Diabetic Normal Normal Overt_Diabetic
## [73] Normal Normal Normal Normal
## [77] Normal Normal Normal Normal
## [81] Overt_Diabetic Normal Overt_Diabetic Overt_Diabetic
## [85] Normal Overt_Diabetic Normal Normal
## [89] Overt_Diabetic Chemical_Diabetic Normal Normal
## [93] Normal Normal Chemical_Diabetic Normal
## [97] Normal Chemical_Diabetic Overt_Diabetic Normal
## [101] Chemical_Diabetic Normal
## Levels: Normal Chemical_Diabetic Overt_Diabetic
```

```
# compute the error rates
```

```
train_error_qda <- (1/length(train_pred_qda$class))*length(which(Y_train != train_pred_qda$class))
test_error_qda <- (1/length(test_pred_qda$class))*length(which(Y_test != test_pred_qda$class))
train_error_qda
```

```
## [1] 0.03921569
```

```
test_error_qda
```

```
## [1] 0.09302326
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:MASS':
##
## select
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
data <- data.frame(Methods = c("LDA" , "QDA"), Train_Error_Percentage = c(train_error_lda * 100,train_error_qda * 100),  
                  Test_Error_Percentage = c(test_error_lda * 100,test_error_qda * 100))  
arrange(data,Test_Error_Percentage)
```

```
##   Methods Train_Error_Percentage Test_Error_Percentage  
## 1    QDA           3.921569           9.302326  
## 2    LDA           9.803922          13.953488
```

From the above error rates clearly we can see that QDA is performing better than LDA.

c) Suppose an individual has (glucose test/intolerance = 68, insulin test =122, SSPG = 544. Relative weight = 1.86, fasting plasma glucose = 184). To which class does LDA assign this individual? To which class does QDA?

```
individual <- data.frame(relwt = 1.86, glufast = 184, glutest = 68, instest = 122, sspg = 544)  
individual
```

```
##   relwt glufast glutest instest sspg  
## 1  1.86   184     68    122  544
```

Predicting using LDA model

```
lda_individual <- predict(lda.fit, newdata = individual)  
lda_individual$class
```

```
## [1] Normal  
## Levels: Normal Chemical_Diabetic Overt_Diabetic
```

Predicting using QDA model

```
qda_individual <- predict(qda.fit, newdata = individual)  
qda_individual$class
```

```
## [1] Overt_Diabetic  
## Levels: Normal Chemical_Diabetic Overt_Diabetic
```

The LDA Model classifies the individual as “Normal” and the QDA Model classifies the individual as “Overt_Diabetic”

Question 2) The insurance company benchmark data set gives information on customers. Specifically, it contains 86 variables on product-usage data and socio-

demographic data derived from zip area codes. There are 5,822 customers in the training set and another 4,000 in the test set. The data were collected to answer the following questions: Can you predict who will be interested in buying a caravan insurance policy and give an explanation why?

```
library(ISLR2)
```

```
##  
## Attaching package: 'ISLR2'
```

```
## The following object is masked from 'package:MASS':  
##  
## Boston
```

```
data("Caravan")  
  
setwd("D:/Buffalo/files")  
  
# Loading the data  
  
train_data <- read.table("Caravan_training_ticdata2000.txt")  
colnames(train_data) <- colnames(Caravan)  
#head(train_data)  
  
X_train <- train_data[,-86]  
  
Y_train <- train_data[,86]  
  
X_test <- read.table("Caravan_test_ticeval2000.txt")  
colnames(X_test) <- colnames(Caravan[,-86])  
#head(X_test)  
  
Y_test <- read.table("Caravan_test_outcome_tictgts2000.txt")  
#head(Y_test)  
table(Y_test)
```

```
## V1  
##    0    1  
## 3762 238
```

```
test_data <- data.frame(X_test,Y_test)  
colnames(test_data) <- colnames(Caravan)  
#head(test_data)
```

```
dim(train_data)
```

```
## [1] 5822   86
```

```
dim(test_data)
```

```
## [1] 4000 86
```

a) Develop a model using the linear model.

```
lm.fit <- lm(Purchase ~. , train_data)  
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Purchase ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.67293 -0.08720 -0.04593 -0.00639  1.04628
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.7685381  0.4298406   1.788 0.073835 .
## MOSTYPE      0.0035209  0.0022512   1.564 0.117866
## MAANTHUI     -0.0072642  0.0076739  -0.947 0.343875
## MGEMOMV      -0.0012739  0.0071737  -0.178 0.859055
## MGEMLEEF     0.0107473  0.0049596   2.167 0.030279 *
## MOSHOOFD    -0.0154869  0.0101044  -1.533 0.125405
## MGODRK       -0.0056016  0.0056016  -1.000 0.317353
## MGODPR       -0.0002069  0.0060664  -0.034 0.972795
## MGODOV       0.0003569  0.0054592   0.065 0.947874
## MGODGE      -0.0030237  0.0058038  -0.521 0.602399
## MRELGE       0.0086829  0.0075479   1.150 0.250036
## MRELSA       0.0020367  0.0072008   0.283 0.777310
## MRELOV       0.0055682  0.0076295   0.730 0.465526
## MFALLEEN    -0.0038250  0.0065474  -0.584 0.559107
## MFG EKIND    -0.0050625  0.0066861  -0.757 0.448980
## MFWEKIND    -0.0026253  0.0069795  -0.376 0.706824
## MOPLHOOG     0.0021357  0.0068161   0.313 0.754038
## MOPLMIDD    -0.0048456  0.0071396  -0.679 0.497358
## MOPLLAAG    -0.0113977  0.0073004  -1.561 0.118525
## MBERHOOG     0.0021884  0.0045182   0.484 0.628153
## MBERZELF    -0.0004665  0.0052201  -0.089 0.928796
## MBERBOER    -0.0050974  0.0050426  -1.011 0.312122
## MBERMIDD     0.0041254  0.0044806   0.921 0.357228
## MBERARBG    -0.0006060  0.0044709  -0.136 0.892190
## MBERARBO     0.0019733  0.0044532   0.443 0.657690
## MSKA        -0.0013674  0.0051653  -0.265 0.791225
## MSKB1       -0.0031701  0.0050198  -0.632 0.527724
## MSKB2       -0.0012603  0.0044827  -0.281 0.778603
## MSKC         0.0024879  0.0049115   0.507 0.612502
## MSKD        -0.0008866  0.0047145  -0.188 0.850832
## MHHUUR      -0.0454201  0.0376622  -1.206 0.227872
## MHKOOP      -0.0432242  0.0376290  -1.149 0.250730
## MAUT1        0.0085964  0.0075592   1.137 0.255502
## MAUT2        0.0077871  0.0068554   1.136 0.256038
## MAUT0        0.0047215  0.0072646   0.650 0.515762
## MZFONDS     -0.0561024  0.0444643  -1.262 0.207094
## MZPART       -0.0593733  0.0443897  -1.338 0.181097
## MINKM30      0.0070879  0.0051150   1.386 0.165884
## MINK3045     0.0069414  0.0049276   1.409 0.158986
## MINK4575     0.0049679  0.0050144   0.991 0.321862
## MINK7512     0.0059267  0.0052728   1.124 0.261053
## MINK123M    -0.0098939  0.0069270  -1.428 0.153258
## MINKGEM      0.0063044  0.0045645   1.381 0.167277
## MKOOPKLA     0.0029097  0.0022664   1.284 0.199250
## PWAPART      0.0284931  0.0166017   1.716 0.086166 .
```

```

## PWABEDR      -0.0101533  0.0205121  -0.495  0.620625
## PWALAND      -0.0201220  0.0390424  -0.515  0.606301
## PPERSAUT      0.0102787  0.0026346   3.901  9.67e-05 ***
## PBESAUT      0.0014405  0.0148574   0.097  0.922765
## PMOTSCO      -0.0061279  0.0079415  -0.772  0.440364
## PVRAAUT      -0.0249190  0.0415892  -0.599  0.549083
## PAANHANG      0.0588044  0.0557610   1.055  0.291662
## PTRACTOR      0.0121481  0.0142358   0.853  0.393504
## PWERKT       -0.0062440  0.0370186  -0.169  0.866060
## PBROM         0.0078683  0.0152793   0.515  0.606598
## PLEVEN       -0.0155397  0.0064753  -2.400  0.016433 *
## PPERSONG      0.0098926  0.0335157   0.295  0.767880
## PGEZONG       0.1937254  0.0793370   2.442  0.014644 *
## PWAOREG       0.0647933  0.0256913   2.522  0.011696 *
## PBRAND        0.0132643  0.0035906   3.694  0.000223 ***
## PZEILPL      -0.1917507  0.1439848  -1.332  0.182998
## PPLEZIER     -0.0299076  0.0269224  -1.111  0.266666
## PFIETS       -0.0107777  0.0549693  -0.196  0.844564
## PINBOED      -0.0441620  0.0307404  -1.437  0.150883
## PBYSTAND     -0.0184858  0.0288890  -0.640  0.522269
## AWAPART      -0.0377952  0.0323794  -1.167  0.243154
## AWABEDR       0.0185448  0.0529740   0.350  0.726296
## AWALAND       0.0180904  0.1374585   0.132  0.895300
## APERSAUT      0.0002821  0.0127496   0.022  0.982347
## ABESAUT      -0.0214816  0.0652955  -0.329  0.742175
## AMOTSCO       0.0203252  0.0310683   0.654  0.513004
## AVRAAUT       0.0563675  0.1589388   0.355  0.722866
## AAANHANG     -0.0804238  0.0944352  -0.852  0.394455
## ATRACTOR     -0.0395651  0.0353795  -1.118  0.263484
## AWERKT       -0.0010526  0.0728240  -0.014  0.988468
## ABROM        -0.0236462  0.0467611  -0.506  0.613101
## ALEVEN        0.0372344  0.0154024   2.417  0.015661 *
## APERSONG     -0.0464279  0.0954471  -0.486  0.626684
## AGEZONG      -0.4050642  0.1898715  -2.133  0.032938 *
## AWAOREG      -0.2304561  0.1243310  -1.854  0.063852 .
## ABRAND       -0.0211374  0.0116048  -1.821  0.068593 .
## AZEILPL       0.4958051  0.2815591   1.761  0.078304 .
## APLEZIER      0.3633887  0.0885318   4.105  4.11e-05 ***
## AFIETS        0.0416061  0.0408644   1.018  0.308650
## AINBOED       0.0959436  0.0699079   1.372  0.169983
## ABYSTAND      0.1312250  0.0983836   1.334  0.182319
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.23 on 5736 degrees of freedom
## Multiple R-squared:  0.0729, Adjusted R-squared:  0.05916
## F-statistic: 5.306 on 85 and 5736 DF,  p-value: < 2.2e-16

```

```

train_predict_lm <- predict(lm.fit, newdata = X_train)
test_predict_lm <- predict(lm.fit, newdata = X_test)

head(train_predict_lm)

```

```
##           1           2           3           4           5           6
## 0.09738541 0.01345938 0.08354523 0.09075754 0.04307400 0.01475749
```

```
head(test_predict_lm)
```

```
##           1           2           3           4           5           6
## 0.014441132 0.215946829 0.099937482 0.095439888 0.005945841 0.027520016
```

```
# since, the target or response is a qualitative variable with 2 classes(0 & 1) we will convert the outcomes to 0 & 1
```

```
train_predict_lm <- ifelse(train_predict_lm > 0.5, 1, 0)
test_predict_lm <- ifelse(test_predict_lm > 0.5, 1, 0)
```

```
train_error_lm <- mean(Y_train != train_predict_lm)
```

```
test_error_lm <- mean(Y_test != test_predict_lm)
```

```
train_error_lm
```

```
## [1] 0.05960151
```

```
test_error_lm
```

```
## [1] 0.05975
```

```
table(Y_test[,],test_predict_lm)
```

```
##      test_predict_lm
##           0           1
## 0 3760           2
## 1  237           1
```

b) Develop a model using Forwards Selection, Backwards Selection, Lasso regression, and Ridge regression.

Forward subset selection:

```
library(leaps)
```

```
# Performing forward subset selection on the data
```

```
regfit.fwd <- regsubsets(Purchase ~., data = train_data, nbest = 1, nvmax = 85, method = "forward")
```

```
my_sum_fwd <- summary(regfit.fwd)
names(my_sum_fwd)
```

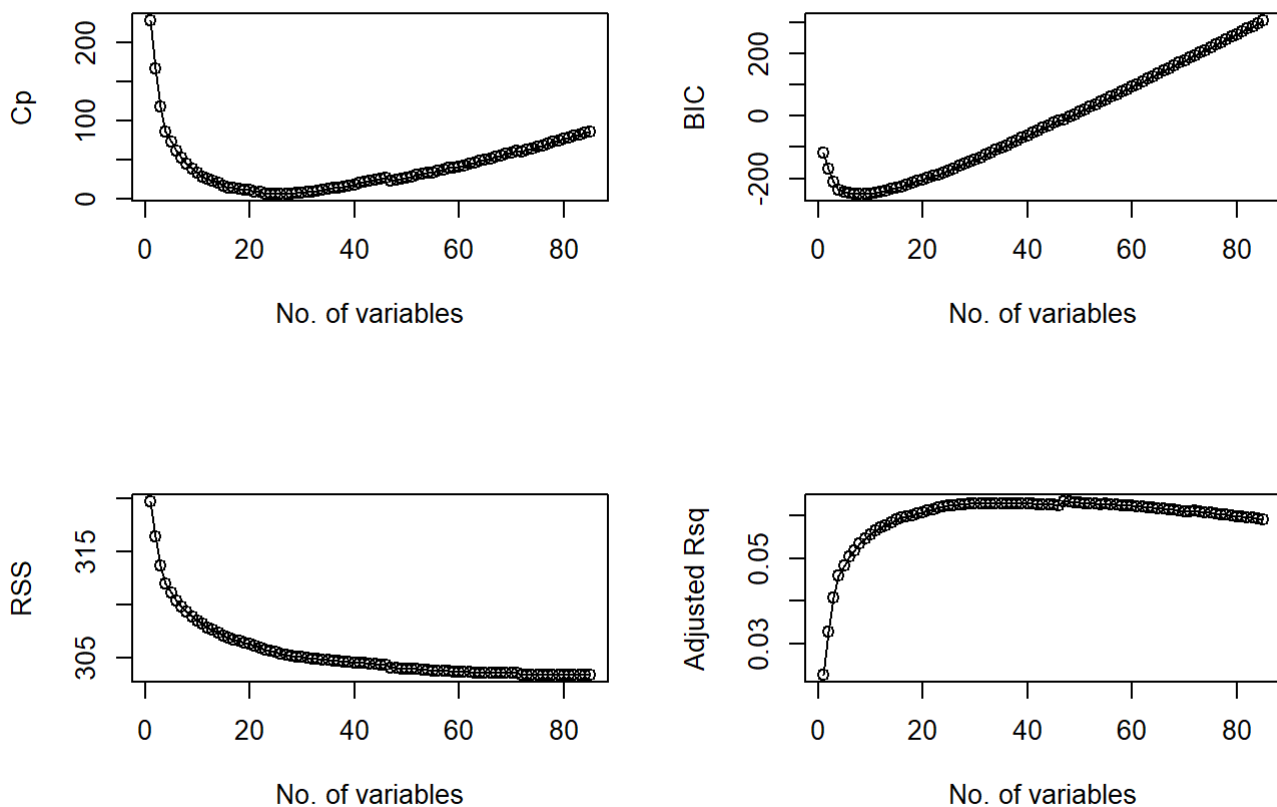
```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
# examine the best "p" variables models
```

```
### my_sum_fwd$outmat
```

```
# plot model selection measures
```

```
par(mfrow = c(2,2))
plot(my_sum_fwd$cp, xlab = "No. of variables", ylab = "Cp", type = "o")
plot(my_sum_fwd$bic, xlab = "No. of variables", ylab = "BIC", type = "o")
plot(my_sum_fwd$rss, xlab = "No. of variables", ylab = "RSS", type = "o")
plot(my_sum_fwd$adjr2, xlab = "No. of variables", ylab = "Adjusted Rsq", type = "o")
```



```
# identify the optimal models using model selection measures for forward subset selection
```

```
data <- data.frame(Parameters = c("CP","BIC","RSS","Adj Rsquare"), Values = c(which.min(my_sum_fwd$cp), which.min(my_sum_fwd$bic), which.min(my_sum_fwd$rss), which.max(my_sum_fwd$adjr2)))
data
```

```
## Parameters Values
## 1 CP 23
## 2 BIC 8
## 3 RSS 85
## 4 Adj Rsquare 47
```



```
### Predicting training and test errors
```

```
predict.regsubsets = function(object, newdata, id){  
  form = as.formula(object$call[[2]])  
  mat = model.matrix(form, newdata)  
  coefi = coef(object,id=id)  
  xvars=names(coefi)  
  mat[,xvars]%*%coefi  
}
```

```
# create objects to store error.
```

```
train_err_store1 <- matrix(rep(NA, 85))
```

```
test_err_store1 <- matrix(rep(NA, 85))
```


```
for (i in 1:85){  
  # make the predictions  
  y_hat_train1 <- predict(regfit.fwd, newdata = train_data, id = i)  
  y_hat_test1 <- predict(regfit.fwd, newdata = test_data, id = i)  
  
  # converting data to 0 & 1  
  
  y_hat_train1 <- ifelse(y_hat_train1 > 0.5, 1, 0)  
  y_hat_test1 <- ifelse(y_hat_test1 > 0.5, 1, 0)  
  
  # compare the prediction with the true  
  train_err_store1[i] <- mean(Y_train != y_hat_train1)  
  test_err_store1[i] <- mean(Y_test != y_hat_test1)  
}
```

```
par(mfrow = c(1,1))
```

```
plot(train_err_store1, col = "blue", type = "b", xlab = "No. of variables", ylab = "MSE",main  
="Forward subset selection MSE")
```

```
lines(test_err_store1, col = "red", type = "b")
```

Forward subset selection MSE



Backward subset selection:

```
regfit.bwd <- regsubsets(Purchase ~., data = train_data, nbest = 1, nvmax = 85, method = "backward")
```

```
my_sum_bwd <- summary(regfit.bwd)
names(my_sum_bwd)
```



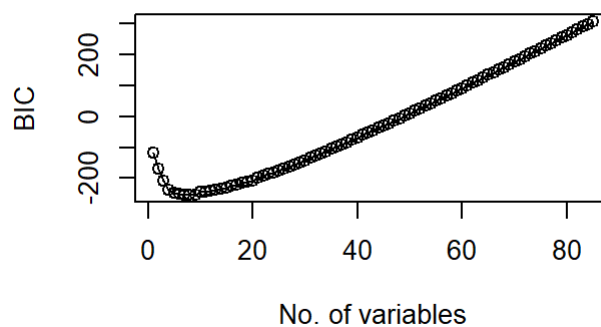
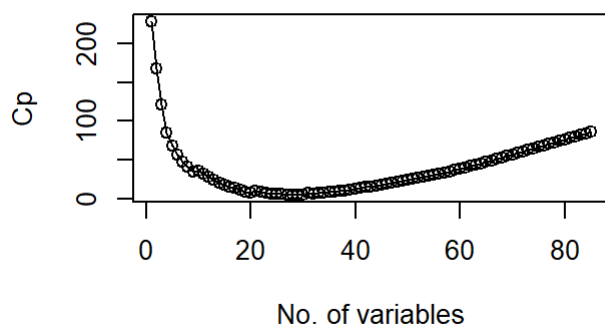
```
## [1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
```

```
# examine the best "p" variables models
```

```
#### my_sum_bwd$outmat
```

```
# plot model selection measures
```

```
par(mfrow = c(2,2))
plot(my_sum_bwd$cp, xlab = "No. of variables", ylab = "Cp", type = "o")
plot(my_sum_bwd$bic, xlab = "No. of variables", ylab = "BIC", type = "o")
plot(my_sum_bwd$rss, xlab = "No. of variables", ylab = "RSS", type = "o")
plot(my_sum_bwd$adjr2, xlab = "No. of variables", ylab = "Adjusted Rsq", type = "o")
```



identify the optimal models using model selection measures for backward subset selection

```
data <- data.frame(Parameters = c("CP","BIC","RSS","Adj Rsquare"), Values = c(which.min(my_sum_bwd$cp), which.min(my_sum_bwd$bic), which.min(my_sum_bwd$rss), which.max(my_sum_bwd$adjr2)))
data
```

##	Parameters	Values
## 1	CP	29
## 2	BIC	8
## 3	RSS	85
## 4	Adj Rsquare	39

```
### Predicting training and test errors
```

```
predict.regsubsets = function(object, newdata, id){  
  form = as.formula(object$call[[2]])  
  mat = model.matrix(form, newdata)  
  coefi = coef(object,id=id)  
  xvars=names(coefi)  
  mat[,xvars]%*%coefi  
}
```

```
# create objects to store error.
```

```
train_err_store2 <- matrix(rep(NA, 85))
```

```
test_err_store2 <- matrix(rep(NA, 85))
```


```
for (i in 1:85){  
  # make the predictions  
  y_hat_train2 <- predict(regfit.bwd, newdata = train_data, id = i)  
  y_hat_test2 <- predict(regfit.bwd, newdata = test_data, id = i)  
  
  # converting data to 0 & 1  
  
  y_hat_train2 <- ifelse(y_hat_train2 > 0.5, 1, 0)  
  y_hat_test2 <- ifelse(y_hat_test2 > 0.5, 1, 0)  
  
  # compare the prediction with the true  
  train_err_store2[i] <- mean(Y_train != y_hat_train2)  
  test_err_store2[i] <- mean(Y_test != y_hat_test2)  
}
```

```
par(mfrow = c(1,1))
```

```
plot(train_err_store2, col = "blue", type = "b", xlab = "No. of variables", ylab = "MSE",main  
="Backward subset selection MSE")
```

```
lines(test_err_store2, col = "red", type = "b")
```

Backward subset selection MSE



Lasso Regression:


```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
set.seed(123)
```

```
lasso_model <- cv.glmnet(x = as.matrix(X_train), y = Y_train, alpha = 1, nfolds = 10)  
plot(lasso_model)
```



```
best_lambda <- lasso_model$lambda.min
best_lambda
```

```
## [1] 0.002644076
```

```
# predicting the response of the variables using best_lambda on training and test data
```

```
train_predict_lr <- predict(lasso_model, s = best_lambda, newx = as.matrix(X_train), type =
'response')
train_predict_lr <- ifelse(train_predict_lr > 0.5, 1, 0)
```

```
test_predict_lr <- predict(lasso_model, s = best_lambda, newx = as.matrix(X_test), type = 're
sponse')
test_predict_lr <- ifelse(test_predict_lr > 0.5, 1, 0)
```

```
# Calculating train and test errors:
```

```
train_error_lr <- mean(Y_train != train_predict_lr)
test_error_lr <- mean(Y_test != test_predict_lr)
```

```
train_error_lr
```

```
## [1] 0.05977327
```

```
test_error_lr
```

```
## [1] 0.05975
```

```
table(Y_test[,],test_predict_lr)
```


```
##      test_predict_lr
##           0         1
##    0 3760         2
##    1  237         1
```

Ridge Regression:

```
library(glmnet)
```

```
set.seed(123)
```

```
ridge_model <- cv.glmnet(x = as.matrix(X_train), y = Y_train, alpha = 0, nfolds = 10)
plot(ridge_model)
```



```
best_lambda <- ridge_model$lambda.min
best_lambda
```

```
## [1] 0.1018902
```

```
# predicting the response of the variables using best_lambda on training and test data

train_predict_rr <- predict(ridge_model, s = best_lambda, newx = as.matrix(X_train), type =
'response')
train_predict_rr <- ifelse(train_predict_lr > 0.5, 1, 0)

test_predict_rr <- predict(ridge_model, s = best_lambda, newx = as.matrix(X_test), type = 're
sponse')
test_predict_rr <- ifelse(test_predict_lr > 0.5, 1, 0)

# Calculating train and test errors:
train_error_rr <- mean(Y_train != train_predict_rr)
test_error_rr <- mean(Y_test != test_predict_rr)

train_error_rr
```

```
## [1] 0.05977327
```

```
test_error_rr
```

```
## [1] 0.05975
```

```
table(Y_test[,], test_predict_rr)
```

```
##      test_predict_rr
##           0      1
##    0 3760      2
##    1  237      1
```

c) Develop a model using logistic regression.

```
library(caret)
```

```
## Loading required package: lattice
```

```
set.seed(123)
glm.fit <- glm(Purchase ~. , data = train_data, family = "binomial")
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm.fit)
```



```
##
## Call:
## glm(formula = Purchase ~ ., family = "binomial", data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  2.542e+02  1.116e+04   0.023   0.98183
## MOSTYPE      6.580e-02  4.624e-02   1.423   0.15468
## MAANTHUI     -1.832e-01  1.927e-01  -0.951   0.34157
## MGEMOMV      -2.696e-02  1.399e-01  -0.193   0.84723
## MGEMLEEF      2.096e-01  1.016e-01   2.063   0.03911 *
## MOSHOOFD     -2.767e-01  2.076e-01  -1.333   0.18247
## MGODRK       -1.142e-01  1.069e-01  -1.068   0.28535
## MGODPR       -1.910e-02  1.177e-01  -0.162   0.87112
## MGODOV       -1.618e-02  1.055e-01  -0.153   0.87818
## MGODGE       -6.817e-02  1.113e-01  -0.612   0.54024
## MRELGE       2.310e-01  1.566e-01   1.475   0.14031
## MRELSA       8.509e-02  1.466e-01   0.580   0.56169
## MRELOV       1.467e-01  1.562e-01   0.939   0.34759
## MFALLEEN     -8.291e-02  1.311e-01  -0.633   0.52702
## MFGEKIND     -1.154e-01  1.337e-01  -0.863   0.38813
## MFWEKIND     -8.140e-02  1.417e-01  -0.575   0.56561
## MOPLHOOG      9.717e-04  1.311e-01   0.007   0.99408
## MOPLMIDD     -9.077e-02  1.365e-01  -0.665   0.50605
## MOPLLAAG     -1.994e-01  1.376e-01  -1.449   0.14740
## MBERHOOG      8.883e-02  9.349e-02   0.950   0.34204
## MBERZELF      3.918e-02  9.897e-02   0.396   0.69219
## MBERBOER     -1.169e-01  1.104e-01  -1.059   0.28951
## MBERMIDD      1.353e-01  9.191e-02   1.472   0.14106
## MBERARBG      3.976e-02  9.067e-02   0.438   0.66104
## MBERARBO      9.954e-02  9.143e-02   1.089   0.27628
## MSKA          2.690e-02  1.035e-01   0.260   0.79502
## MSKB1        -8.801e-03  1.011e-01  -0.087   0.93064
## MSKB2         1.200e-02  9.081e-02   0.132   0.89485
## MSKC          9.016e-02  9.958e-02   0.905   0.36527
## MSKD         -2.468e-02  9.724e-02  -0.254   0.79967
## MHHUUR       -1.472e+01  8.140e+02  -0.018   0.98557
## MHKOOP       -1.469e+01  8.140e+02  -0.018   0.98561
## MAUT1         1.819e-01  1.514e-01   1.202   0.22953
## MAUT2         1.507e-01  1.371e-01   1.099   0.27162
## MAUT0         9.325e-02  1.436e-01   0.649   0.51603
## MZFONDS      -1.445e+01  9.359e+02  -0.015   0.98768
## MZPART       -1.451e+01  9.359e+02  -0.016   0.98763
## MINKM30       1.181e-01  1.006e-01   1.174   0.24039
## MINK3045      1.366e-01  9.650e-02   1.415   0.15694
## MINK4575      1.009e-01  9.667e-02   1.043   0.29678
## MINK7512      1.144e-01  1.027e-01   1.114   0.26513
## MINK123M     -1.607e-01  1.449e-01  -1.109   0.26738
## MINKGEM       9.214e-02  9.945e-02   0.927   0.35417
## MKOOPKLA      6.856e-02  4.642e-02   1.477   0.13966
## PWAPART       5.954e-01  3.901e-01   1.526   0.12693
## PWABEDR      -2.757e-01  4.635e-01  -0.595   0.55196
## PWALAND      -4.405e-01  1.035e+00  -0.425   0.67052
## PERSAUT       2.306e-01  4.199e-02   5.491  4.01e-08 ***
## PBESAUT       1.215e+01  4.029e+02   0.030   0.97595
```

```

## PMOTSCO      -8.101e-02  1.147e-01  -0.706  0.48006
## PVRAAUT      -2.106e+00  2.557e+03  -0.001  0.99934
## PAANHANG      1.014e+00  9.371e-01   1.082  0.27917
## PTRACTOR      7.229e-01  4.278e-01   1.690  0.09107 .
## PWERKT      -5.525e+00  4.805e+03  -0.001  0.99908
## PBROM        2.170e-01  4.865e-01   0.446  0.65559
## PLEVEN       -2.382e-01  1.170e-01  -2.036  0.04173 *
## PPERSONG     -4.523e-01  2.094e+00  -0.216  0.82901
## PGEZONG      1.444e+00  1.029e+00   1.404  0.16033
## PWAOREG      8.239e-01  5.943e-01   1.386  0.16565
## PBRAND       2.401e-01  7.714e-02   3.113  0.00185 **
## PZEILPL     -8.658e+00  3.261e+03  -0.003  0.99788
## PPLEZIER     -1.886e-01  3.259e-01  -0.579  0.56289
## PFIETS       3.664e-01  8.325e-01   0.440  0.65985
## PINBOED     -1.068e+00  8.764e-01  -1.219  0.22301
## PBYSTAND     -1.676e-01  3.321e-01  -0.505  0.61373
## AWAPART     -9.293e-01  7.802e-01  -1.191  0.23364
## AWABEDR      4.197e-01  1.082e+00   0.388  0.69824
## AWALAND      2.762e-01  3.528e+00   0.078  0.93758
## APERSAUT     -3.902e-02  1.772e-01  -0.220  0.82566
## ABESAUT     -7.298e+01  2.417e+03  -0.030  0.97591
## AMOTSCO      2.418e-01  3.772e-01   0.641  0.52142
## AVRAAUT     -4.490e+00  1.078e+04   0.000  0.99967
## AAANHANG     -1.351e+00  1.687e+00  -0.801  0.42322
## ATRACTOR     -2.376e+00  1.524e+00  -1.559  0.11899
## AWERKT      -8.749e-01  9.682e+03   0.000  0.99993
## ABROM       -1.060e+00  1.549e+00  -0.684  0.49367
## ALEVEN       4.789e-01  2.245e-01   2.133  0.03291 *
## APERSONG      3.997e-01  4.329e+00   0.092  0.92644
## AGEZONG     -3.163e+00  2.706e+00  -1.169  0.24247
## AWAOREG     -3.212e+00  3.433e+00  -0.936  0.34939
## ABRAND       -4.118e-01  2.787e-01  -1.477  0.13956
## AZEILPL      1.047e+01  3.261e+03   0.003  0.99744
## APLEZIER      2.516e+00  1.010e+00   2.490  0.01276 *
## AFIETS       2.318e-01  5.699e-01   0.407  0.68420
## AINBOED      1.947e+00  1.412e+00   1.378  0.16812
## ABYSTAND      1.078e+00  1.103e+00   0.977  0.32870
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2635.5  on 5821  degrees of freedom
## Residual deviance: 2243.5  on 5736  degrees of freedom
## AIC: 2415.5
##
## Number of Fisher Scoring iterations: 17

```

```
names(glm.fit)
```

```
## [1] "coefficients"      "residuals"      "fitted.values"
## [4] "effects"           "R"              "rank"
## [7] "qr"                "family"         "linear.predictors"
## [10] "deviance"          "aic"            "null.deviance"
## [13] "iter"              "weights"        "prior.weights"
## [16] "df.residual"       "df.null"        "y"
## [19] "converged"         "boundary"       "model"
## [22] "call"              "formula"        "terms"
## [25] "data"              "offset"         "control"
## [28] "method"            "contrasts"      "xlevels"
```

```
# Predictions
```

```
glm.probs.train <- predict(glm.fit, newdata = train_data, type = "response")
y_hat_train <- ifelse(glm.probs.train > 0.5, 1, 0)

glm.probs.test <- predict(glm.fit, newdata = test_data, type = "response")
y_hat_test <- ifelse(glm.probs.test > 0.5, 1, 0)
```

```
# Calculate the error rates
```

```
train_err <- mean(y_hat_train != Y_train)
train_err
```

```
## [1] 0.05994504
```

```
test_err <- mean(y_hat_test != Y_test)
test_err
```

```
## [1] 0.06025
```

```
# Confusion matrix
```

```
conf <- confusionMatrix(as.factor(y_hat_test), as.factor(Y_test[ , ]))
conf$table
```

```
##           Reference
## Prediction    0    1
##           0 3756  235
##           1    6    3
```

The MSE of all the methods forward subset selection, backward subset selection, Lasso Regression, Ridge Regression and logistic regression are very close.

Lasso and Ridge have similar test accuracy.

And our models didn't performed well in predicting Yes(1's), The reason might be number of 1's are very less in the given data.

Question 3) In this exercise, we will predict the number of applications received using the other variables in the College data set.

```
library(caret)
library(ISLR2)
library(glmnet)
library(pls)
```

```
##
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:caret':
##
##      R2
```

```
## The following object is masked from 'package:corrplot':
##
##      corrplot
```

```
## The following object is masked from 'package:stats':
##
##      loadings
```

```
data("College")
dim(College) ## 777 * 18
```

```
## [1] 777  18
```

```
head(College)
```

##	Private	Apps	Accept	Enroll	Top10perc	Top25perc
## Abilene Christian University	Yes	1660	1232	721	23	52
## Adelphi University	Yes	2186	1924	512	16	29
## Adrian College	Yes	1428	1097	336	22	50
## Agnes Scott College	Yes	417	349	137	60	89
## Alaska Pacific University	Yes	193	146	55	16	44
## Albertson College	Yes	587	479	158	38	62
##	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	
## Abilene Christian University	2885	537	7440	3300	450	
## Adelphi University	2683	1227	12280	6450	750	
## Adrian College	1036	99	11250	3750	400	
## Agnes Scott College	510	63	12960	5450	450	
## Alaska Pacific University	249	869	7560	4120	800	
## Albertson College	678	41	13500	3335	500	
##	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend
## Abilene Christian University	2200	70	78	18.1	12	7041
## Adelphi University	1500	29	30	12.2	16	10527
## Adrian College	1165	53	66	12.9	30	8735
## Agnes Scott College	875	92	97	7.7	37	19016
## Alaska Pacific University	1500	76	72	11.9	2	10922
## Albertson College	675	67	73	9.4	11	9727
##	Grad.Rate					
## Abilene Christian University	60					
## Adelphi University	56					
## Adrian College	54					
## Agnes Scott College	59					
## Alaska Pacific University	15					
## Albertson College	55					

```
### Private variable is a factor changing it to 1(Yes) and 0(No)
College$Private <- ifelse(College$Private == 'Yes', 1 , 0)
```

a) Split the data set into a training set and a test set.

```
set.seed(123)
train_indis <- sample(c(1:length(College[,1])), size = round(2/3*length(College[,1])), replace = FALSE)

train_data <- College[train_indis, ]
test_data <- College[-train_indis, ]

X_train <- train_data[, -2]
Y_train <- train_data[, 2]

X_test <- test_data[, -2]
Y_test <- test_data[, 2]
```

b) Fit a linear model using least squares on the training set, and report the test error obtained.

```
set.seed(123)
lm.fit <- lm(Apps ~. , data = train_data)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Apps ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3098.1  -435.7   -32.6    326.9   6524.3
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -320.63000   483.82540   -0.663  0.507830
## Private      -631.06608   166.38884   -3.793  0.000167 ***
## Accept         1.22765     0.05907   20.782 < 2e-16 ***
## Enroll         0.07342     0.22242    0.330  0.741483
## Top10perc     45.28449     6.30692    7.180  2.54e-12 ***
## Top25perc    -12.88783     5.12008   -2.517  0.012144 *
## F.Undergrad    0.02496     0.04024    0.620  0.535386
## P.Undergrad    0.03394     0.03505    0.968  0.333304
## Outstate     -0.06350     0.02155   -2.947  0.003361 **
## Room.Board     0.20100     0.05392    3.728  0.000215 ***
## Books          0.16346     0.27890    0.586  0.558084
## Personal      -0.03987     0.07418   -0.537  0.591204
## PhD           -6.76818     5.36695   -1.261  0.207866
## Terminal      -5.29390     5.82889   -0.908  0.364201
## S.F.Ratio     -0.13458    14.77294   -0.009  0.992735
## perc.alumni   -7.16431     4.68079   -1.531  0.126506
## Expend         0.08032     0.01338    6.005  3.69e-09 ***
## Grad.Rate      9.82319     3.37117    2.914  0.003730 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 980.1 on 500 degrees of freedom
## Multiple R-squared:  0.918, Adjusted R-squared:  0.9153
## F-statistic: 329.5 on 17 and 500 DF,  p-value: < 2.2e-16
```

```
ols_predict <- predict(lm.fit, newdata = test_data)
ols_mse <- mean((Y_test - ols_predict)^2)
ols_mse
```


```
## [1] 1684049
```

c) Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
set.seed(123)
ridge_cv <- cv.glmnet(as.matrix(X_train), Y_train, alpha = 0, nfolds = 10)
sel <- ridge_cv$lambda.min
sel
```

```
## [1] 311.779
```

```
plot(ridge_cv)
```



```
ridge_predict <- predict(ridge_cv, s = sel, newx= as.matrix(X_test))
ridge_mse <- (1/length(Y_test))*(sum((ridge_predict - Y_test)^2))
ridge_mse
```


```
## [1] 2787195
```

d) Fit a lasso model on the training set, with λ chosen by crossvalidation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```
set.seed(123)
lasso_cv <- cv.glmnet(as.matrix(X_train), Y_train, alpha = 1, nfolds = 10)
sel <- lasso_cv$lambda.min
sel
```

```
## [1] 6.120348
```

```
plot(lasso_cv)
```



```
lasso_coef <- predict(lasso_cv, s = sel, type = "coefficients")

lasso_predict <- predict(lasso_cv, s = sel, newx= as.matrix(X_test))
lasso_mse <- (1/length(Y_test))*(sum((lasso_predict - Y_test)^2))
lasso_mse
```

```
## [1] 1685197
```

```
lasso_coef
```



```
## 18 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept) -407.16052726
## Private      -612.48955542
## Accept       1.22519081
## Enroll       0.07885531
## Top10perc    41.58235489
## Top25perc    -9.81499052
## F.Undergrad  0.02513096
## P.Undergrad  0.02776503
## Outstate     -0.05735665
## Room.Board   0.18884357
## Books        0.11837682
## Personal     -0.02333079
## PhD          -5.99235815
## Terminal     -5.16787774
## S.F.Ratio     .
## perc.alumni  -6.79379825
## Expend       0.07920785
## Grad.Rate    8.84826047
```

```
number_of_nonzero_coefficients <- sum(lasso_coef[-1] != 0) ## Exclude intercept
number_of_nonzero_coefficients
```

```
## [1] 16
```


e) Fit a PCR model on the training set, with M chosen by cross validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
set.seed(123)
pcr.fit <- pcr(Apps ~ ., data = train_data , scale = TRUE , validation = "CV")
summary(pcr.fit)
```

```
## Data:      X dimension: 518 17
## Y dimension: 518 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           3370    3336    1680    1631    1363    1303    1257
## adjCV        3370    3336    1678    1630    1357    1299    1253
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1202    1201    1169    1169    1168    1174    1176
## adjCV        1195    1196    1166    1167    1165    1171    1173
##      14 comps 15 comps 16 comps 17 comps
## CV           1176    1176    1029    1029
## adjCV        1173    1173    1025    1025
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          31.765   57.84   64.68   70.19   75.49   80.39   84.01   87.40
## Apps       3.386   75.80   77.45   84.75   86.02   86.91   88.03   88.22
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          90.57   93.02   95.07   96.93   98.02   98.88   99.40
## Apps       88.84   88.89   88.94   88.98   89.03   89.03   89.23
##      16 comps 17 comps
## X          99.82   100.0
## Apps       91.74   91.8
```

```
validationplot (pcr.fit , val.type = "MSEP")
```

Apps



```
### From the graph we can clearly see that cross validation selected  $M = P = 16$ 
```

```
### Evaluating the test MSE
```

```
pcr_predict <- predict(pcr.fit,X_test, ncomp=16)
pcr_mse <- mean((pcr_predict - Y_test)^2)
pcr_mse
```

```
## [1] 1785303
```

f) Fit a PLS model on the training set, with M chosen by cross validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
set.seed(123)
pls.fit <- plsr(Apps ~ ., data = train_data , scale = TRUE , validation = "CV")
summary (pls.fit)
```


```

## Data:      X dimension: 518 17
## Y dimension: 518 1
## Fit method: kernelpls
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              3370    1513    1233    1138    1121    1099    1045
## adjCV           3370    1511    1236    1136    1117    1092    1040
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1031    1028    1029    1030    1027    1028    1028
## adjCV        1027    1025    1026    1026    1024    1025    1025
##      14 comps 15 comps 16 comps 17 comps
## CV           1029    1029    1029    1029
## adjCV        1025    1025    1025    1025
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          26.30   42.01   63.26   67.75   71.41   74.08   77.53   80.83
## Apps       80.53   86.92   89.34   90.16   91.05   91.71   91.77   91.79
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X          83.35   86.14   89.53   91.21   93.22   94.67   97.06
## Apps       91.79   91.80   91.80   91.80   91.80   91.80   91.80
##      16 comps 17 comps
## X          99.11   100.0
## Apps       91.80   91.8

```

```
validationplot (pls.fit , val.type = "MSEP")
```

Apps



```
### Cross-validation selected M = 11 as the number of principal components to minimize the out-of-sample MSE.
pls_predict <- predict(pls.fit,X_test,ncomp=11)
pls_mse <- mean((pls_predict - Y_test)^2)
pls_mse
```

```
## [1] 1703326
```

g) Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?

$$R^2 = 1 - \frac{\text{Residual sum of squares (RSS)}}{\text{total sum of squares (TSS)}},$$

$$= 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}.$$

```
library(dplyr)

TSS <- sum((Y_test - mean(Y_test))^2)

data <- data.frame(method = c("OLS", "Ridge", "Lasso", "PCR", "PLS"),
  test_MSE = c(ols_mse, ridge_mse, lasso_mse, pcr_mse, pls_mse),
  test_R2 = c(1 - sum((Y_test - ols_predict)^2) / TSS,
    1 - sum((Y_test - ridge_predict)^2) / TSS,
    1 - sum((Y_test - lasso_predict)^2) / TSS,
    1 - sum((Y_test - pcr_predict)^2) / TSS,
    1 - sum((Y_test - pls_predict)^2) / TSS))
arrange(data, desc(test_R2))
```

```
##   method test_MSE   test_R2
## 1    OLS 1684049 0.9240638
## 2   Lasso 1685197 0.9240120
## 3    PLS 1703326 0.9231946
## 4    PCR 1785303 0.9194981
## 5   Ridge 2787195 0.8743213
```

Each of the models performed well, with test R2 values above 0.91 except Ridge Regression. There is really a minimal performance difference between 4 methods.

The only model that could potentially have a meaningful performance difference was ridge regression (which performed the worst).

Question 4) We will now try to predict per capita crime rate in the Boston data set.

```
library(leaps)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(MASS)
data("Boston")
dim(Boston)  # 506 * 14
```

```
## [1] 506  14
```

```
str(Boston)
```

```
## 'data.frame':  506 obs. of  14 variables:
## $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn     : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas   : int   0 0 0 0 0 0 0 0 0 0 ...
## $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm     : num  6.58 6.42 7.18 7 7.15 ...
## $ age    : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis    : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad    : int   1 2 2 3 3 3 5 5 5 5 ...
## $ tax    : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black  : num  397 397 393 395 397 ...
## $ lstat  : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv   : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
summary(Boston)
```

```

##      crim              zn          indus          chas
## Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   Min.   :0.00000
## 1st Qu.: 0.08205   1st Qu.: 0.00   1st Qu.: 5.19   1st Qu.:0.00000
## Median : 0.25651   Median : 0.00   Median : 9.69   Median :0.00000
## Mean   : 3.61352   Mean   : 11.36   Mean   :11.14   Mean   :0.06917
## 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10   3rd Qu.:0.00000
## Max.   :88.97620   Max.   :100.00   Max.   :27.74   Max.   :1.00000
##      nox              rm          age          dis
## Min.   :0.3850   Min.   :3.561   Min.   : 2.90   Min.   : 1.130
## 1st Qu.:0.4490   1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.100
## Median :0.5380   Median :6.208   Median : 77.50   Median : 3.207
## Mean   :0.5547   Mean   :6.285   Mean   : 68.57   Mean   : 3.795
## 3rd Qu.:0.6240   3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.188
## Max.   :0.8710   Max.   :8.780   Max.   :100.00   Max.   :12.127
##      rad              tax          ptratio      black
## Min.   : 1.000   Min.   :187.0   Min.   :12.60   Min.   : 0.32
## 1st Qu.: 4.000   1st Qu.:279.0   1st Qu.:17.40   1st Qu.:375.38
## Median : 5.000   Median :330.0   Median :19.05   Median :391.44
## Mean   : 9.549   Mean   :408.2   Mean   :18.46   Mean   :356.67
## 3rd Qu.:24.000   3rd Qu.:666.0   3rd Qu.:20.20   3rd Qu.:396.23
## Max.   :24.000   Max.   :711.0   Max.   :22.00   Max.   :396.90
##      lstat          medv
## Min.   : 1.73   Min.   : 5.00
## 1st Qu.: 6.95   1st Qu.:17.02
## Median :11.36   Median :21.20
## Mean   :12.65   Mean   :22.53
## 3rd Qu.:16.95   3rd Qu.:25.00
## Max.   :37.97   Max.   :50.00

```

(a) Try out some of the regression methods explored in this chapter, such as best subset selection, the lasso, ridge regression, and PCR. Present and discuss results for the approaches that you consider.


```
### best subset selection:
```

```
predict.regsubsets <- function(object, newdata, id){  
  form <- as.formula(object$call[[2]])  
  mat <- model.matrix(form, newdata)  
  coefi <- coef(object,id=id)  
  xvars <- names(coefi)  
  mat[,xvars] %*% coefi  
}
```

```
## For K=10 folds
```


```
k <- 10  
n <- nrow (Boston)  
set.seed (1)  
folds <- sample(rep(1:k, length = n))  
cv.errors <- matrix (NA, k, 13, dimnames = list(NULL, paste(1:13)))  
  
for (i in 1:k) {  
  best.fit <- regsubsets(crim ~ ., data = Boston[folds != i, ], nvmax = 13)  
  
  for (j in 1:13) {  
    pred <- predict(best.fit , Boston[folds == i, ], id = j)  
    cv.errors[i, j] <- mean((Boston$crim[folds == i] - pred)^2)  
  }  
}  
  
mse.cv <- apply (cv.errors , 2, mean)  
mse.cv
```

```
##      1      2      3      4      5      6      7      8  
## 46.00617 44.22854 43.80757 44.63674 44.37501 44.10329 43.38296 43.18012  
##      9     10     11     12     13  
## 42.81453 43.01895 43.03912 42.88730 42.92625
```

```
paste("CV MSE: " , mse.cv[which.min(mse.cv)])
```

```
## [1] "CV MSE: 42.8145280588506"
```

```
par (mfrow = c(1, 1))  
plot (mse.cv , type = "o",pch = 19, col = 'blue', xlab = 'Subset size', ylab = 'CV MSE')
```



```
which.min(mse.cv)
```


```
## 9  
## 9
```

```
### Best subset model selecting 9 parameters or variables.
```

```
### Lasso regression:
```

```
set.seed(1)
```

```
x <- model.matrix(crim ~ . - 1, data = Boston)  
y <- Boston$crim  
cv.lasso <- cv.glmnet(x, y, alpha = 1, type.measure = "mse")  
plot(cv.lasso)
```



```
coef(cv.lasso)
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 1.0894283
## zn          .
## indus       .
## chas        .
## nox         .
## rm          .
## age         .
## dis         .
## rad         0.2643196
## tax         .
## ptratio     .
## black       .
## lstat       .
## medv        .
```

```
cv.lasso$cvm[cv.lasso$lambda == cv.lasso$lambda.1se]
```

```
## [1] 55.3338
```

```
### Ridge regression:
```


```
set.seed(1)
```

```
x <- model.matrix(crim ~ . - 1, data = Boston)
```

```
y <- Boston$crim
```

```
cv.ridge <- cv.glmnet(x, y, alpha = 0, type.measure = "mse")
```

```
plot(cv.ridge)
```



```
coef(cv.ridge)
```

```
## 14 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              s1
```

```
## (Intercept)  1.017516864
```

```
## zn          -0.002805664
```

```
## indus       0.034405928
```

```
## chas       -0.225250602
```

```
## nox        2.249887499
```

```
## rm        -0.162546004
```

```
## age        0.007343331
```

```
## dis       -0.114928730
```

```
## rad        0.059813844
```

```
## tax        0.002659110
```

```
## ptratio    0.086423005
```

```
## black     -0.003342067
```

```
## lstat      0.044495213
```

```
## medv     -0.029124577
```

```
cv.ridge$cvm[cv.ridge$lambda == cv.ridge$lambda.1se]
```

```
## [1] 55.80448
```

```
### Principal Component Regression
```

```
library(pls)
```

```
##  
## Attaching package: 'pls'
```


```
## The following object is masked from 'package:stats':  
##  
## loadings
```

```
set.seed (1)
```

```
pcr.fit <- pcr(crim ~ ., data = Boston , scale = TRUE , validation = "CV")  
summary(pcr.fit)
```

```
## Data:      X dimension: 506 13  
## Y dimension: 506 1  
## Fit method: svdpc  
## Number of components considered: 13  
##  
## VALIDATION: RMSEP  
## Cross-validated using 10 random segments.  
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  
## CV              8.61    7.250    7.253    6.833    6.815    6.826    6.847  
## adjCV           8.61    7.245    7.247    6.825    6.803    6.818    6.838  
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps  
## CV          6.837    6.710    6.735    6.723    6.714    6.696    6.624  
## adjCV       6.827    6.698    6.724    6.710    6.702    6.682    6.609  
##  
## TRAINING: % variance explained  
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps  
## X          47.70    60.36    69.67    76.45    82.99    88.00    91.14    93.45  
## crim       30.69    30.87    39.27    39.61    39.61    39.86    40.14    42.47  
##      9 comps 10 comps 11 comps 12 comps 13 comps  
## X          95.40    97.04    98.46    99.52    100.0  
## crim       42.55    42.78    43.04    44.13    45.4
```

```
validationplot(pcr.fit , val.type = "MSEP")
```



```
### pcr() function reports the root mean squared error. 13 component analysis has the lowest RMSE = 6.624
```

b) Propose a model (or set of models) that seem to perform well on this data set, and justify your answer. Make sure that you are evaluating model performance using validation set error, cross validation, or some other reasonable alternative, as opposed to using training error.

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:MASS':  
##  
##   select
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
data <- data.frame(Methods = c("Best subset selection", "Lasso Regression", "Ridge Regression", "Principal Component Reg"),
  MSE = c(mse.cv[which.min(mse.cv)], cv.lasso$cvm[cv.lasso$lambda == cv.lasso$lambda.1se], cv.ridge$cvm[cv.ridge$lambda == cv.ridge$lambda.1se], 43.877376),
  RMSE = c(sqrt(mse.cv[which.min(mse.cv)]), sqrt(cv.lasso$cvm[cv.lasso$lambda == cv.lasso$lambda.1se]), sqrt(cv.ridge$cvm[cv.ridge$lambda == cv.ridge$lambda.1se]), 6.624)
)
arrange(data, RMSE)
```

```
##           Methods      MSE      RMSE
## 1 Best subset selection 42.81453 6.543281
## 2 Principal Component Reg 43.87738 6.624000
## 3      Lasso Regression 55.33380 7.438669
## 4      Ridge Regression 55.80448 7.470240
```

As i evaluated the performance of the models using cross-validation in part 'a' of this question, Best subset selection model performed better compared to other models with MSE = 42.81 and RMSE = 6.54

c) Does your chosen model involve all of the features in the data set? Why or why not?

```
### The model best subset selection chooses 9 features of the data set. Because, we are getting the minimum MSE = 42.81 at subset size = 9.
```

```
which.min(mse.cv)
```

```
## 9
## 9
```

```
coef(best.fit, id = 9)
```

```
##      (Intercept)           zn           nox           rm           dis
## 12.021667175    0.039587625 -11.310819327    0.831219883   -0.773294339
##           rad      ptratio           black           lstat           medv
## 0.525470938   -0.295718930   -0.006769892    0.179943832   -0.187327063
```

```
### As far as - adding additional features will reduce the training RSS, but causes the cross-validation MSE to increase.
```