

Homework 5

Naga Kartheek Peddisetty, 50538422

12/10/2023

Question 3) Fit a single layer neural network to Weekly data in the ISLR2 package. Predict the direction (Up or Down) of return on a given week. Use cross-validation or the hold out method to determine the number of neurons to use in the layer. Compare your results to those for the logistic regression model. When making the comparison, consider both the classification performance and interpretability of the final model.

```
library(neuralnet)
library(nnet)
library(ISLR2)

data("Weekly")
dim(Weekly)
```

```
## [1] 1089    9
```

```
head(Weekly)
```

```
##   Year  Lag1  Lag2  Lag3  Lag4  Lag5  Volume  Today Direction
## 1 1990  0.816  1.572 -3.936 -0.229 -3.484 0.1549760 -0.270      Down
## 2 1990 -0.270  0.816  1.572 -3.936 -0.229 0.1485740 -2.576      Down
## 3 1990 -2.576 -0.270  0.816  1.572 -3.936 0.1598375  3.514       Up
## 4 1990  3.514 -2.576 -0.270  0.816  1.572 0.1616300  0.712       Up
## 5 1990  0.712  3.514 -2.576 -0.270  0.816 0.1537280  1.178       Up
## 6 1990  1.178  0.712  3.514 -2.576 -0.270 0.1544440 -1.372      Down
```

```
table(Weekly$Direction)
```

```
##  
## Down   Up  
##  484  605
```

```
weekly <- Weekly  
  
weekly$Direction <- as.numeric(weekly$Direction) - 1  
  
head(weekly)
```

```
##   Year  Lag1  Lag2  Lag3  Lag4  Lag5  Volume  Today Direction  
## 1 1990  0.816  1.572 -3.936 -0.229 -3.484 0.1549760 -0.270      0  
## 2 1990 -0.270  0.816  1.572 -3.936 -0.229 0.1485740 -2.576      0  
## 3 1990 -2.576 -0.270  0.816  1.572 -3.936 0.1598375  3.514      1  
## 4 1990  3.514 -2.576 -0.270  0.816  1.572 0.1616300  0.712      1  
## 5 1990  0.712  3.514 -2.576 -0.270  0.816 0.1537280  1.178      1  
## 6 1990  1.178  0.712  3.514 -2.576 -0.270 0.1544440 -1.372      0
```

```
table(weekly$Direction)
```

```
##  
##    0    1  
## 484 605
```

```
set.seed(123)

indis <- sample(1:nrow(weekly),size = round(0.65 * nrow(weekly)), replace = FALSE)

train_data <- weekly[indis, ]
test_data <- weekly[-indis, ]

X_train <- train_data[, -9]
Y_train <- train_data[, 9]

X_test <- test_data[, -9]
Y_test <- test_data[, 9]

train_err_store <- c()
test_err_store <- c()

for (i in 1:10){
  # fitting neural network with "i" neurons
  nn <- neuralnet(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = train_data, hidden = i, err.fct = "ce", st
  epmax = 10^9, linear.output = FALSE)

  pred_train <- predict(nn, newdata = train_data)
  y_hat_train <- ifelse(pred_train > 0.5 , 1 , 0)
  train_err <- sum(Y_train != y_hat_train) / nrow(train_data)
  train_err_store <- c(train_err_store, train_err)

  pred_test <- predict(nn, newdata = test_data)
  y_hat_test <- ifelse(pred_test > 0.5, 1, 0)
  test_err <- sum(Y_test != y_hat_test) / nrow(test_data)
  test_err_store <- c(test_err_store, test_err)

  if (i == 5) {
    plot(nn)
    print(table(Actual = Y_test, Predicted = y_hat_test))
  }
}
```

```
##      Predicted
## Actual    0    1
##         0 66 92
##         1 80 143
```

```
train_err_store
```

```
## [1] 0.4124294 0.4053672 0.3855932 0.3559322 0.3403955 0.3262712 0.2711864
## [8] 0.3036723 0.2923729 0.2923729
```

```
test_err_store
```

```
## [1] 0.4829396 0.4619423 0.4566929 0.4619423 0.4514436 0.4724409 0.4750656
## [8] 0.4566929 0.4803150 0.4829396
```

```
test_err_store[which.min(test_err_store)]
```

```
## [1] 0.4514436
```

```
paste(which.min(test_err_store),"Neurons optimal single layer neural network model with minimum test error : " , round(test_err_store[which.min(test_err_store)],3))
```

```
## [1] "5 Neurons optimal single layer neural network model with minimum test error : 0.451"
```

Logistic Regression

```
set.seed(123)

glm.fit <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, family = "binomial", data = train_data)

pred_test <- predict(glm.fit , newdata = test_data, type="response")
y_hat_test <- ifelse(pred_test > 0.5, 1, 0)

test_err_glm <- sum(y_hat_test != Y_test) / nrow(test_data)
table(Actual = Y_test, Predicted = y_hat_test)
```

```
##      Predicted
## Actual    0    1
##      0  32 126
##      1   35 188
```

```
test_err_glm
```

```
## [1] 0.4225722
```

Comparision between Neural Network and Logistic Regression

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:neuralnet':
##
##      compute
```

```
## The following objects are masked from 'package:stats':
##
##      filter, lag
```

```
## The following objects are masked from 'package:base':
##
## intersect, setdiff, setequal, union
```

```
a <- data.frame(Methods = c("single layer neural network model with five neurons", "Logistic Regression"), Test_MSE = c(test_err_store[which.min(test_err_store)], test_err_glm))

arrange(a, Test_MSE)
```

```
##
##           Methods Test_MSE
## 1           Logistic Regression 0.4225722
## 2 single layer neural network model with five neurons 0.4514436
```

On the Weekly data logistic regression is performing better than single layer neural network when compared with test_mse.

The interpretability of logistic regression is also simple when compared to the neural networks even with single layer has more complex model and difficult to interpret the weights and biases.

Question 4) Take any classification data set and divide it up into a learning set and a test set. Change the value of one observation on one input variable in the learning set so that the value is now a univariate outlier. Fit separate single-hidden layer neural networks to the original learning-set data and to the learning-set data with the outlier. Use cross-validation or the hold out method to determine the number of neurons to use in the layer. Comment on the effect of the outlier on the fit and on its effect on classifying the test set. Shrink the value of that outlier toward its original value and evaluate when the effect of the outlier on the fit vanishes. How far away must the outlier move from its original value that significant changes to the network coefficient estimates occur?

```
library(neuralnet)
library(nnet)

data(infert)
dim(infert) ### 248 * 8
```

```
## [1] 248 8
```

```
head(infert)
```

```
##   education age parity induced case spontaneous stratum pooled.stratum
## 1    0-5yrs  26     6      1    1           2        1           3
## 2    0-5yrs  42     1      1    1           0        2           1
## 3    0-5yrs  39     6      2    1           0        3           4
## 4    0-5yrs  34     4      2    1           0        4           2
## 5    6-11yrs 35     3      1    1           1        5          32
## 6    6-11yrs 36     4      2    1           1        6          36
```

```
table(infert$case)
```

```
##
##   0    1
## 165  83
```

```
inf <- infert

set.seed(123)

indis <- sample(1:nrow(inf), round(2/3*nrow(inf)), replace = FALSE)
train <- inf[indis, ]
test <- inf[-indis, ]

### Adding outlier to age variable in the first row
train_outlier <- train
train_outlier$age[1] <- 100

train_err_store <- c()
test_err_store <- c()
for (i in 1:5){

  # fit neural network with "i" neurons
  nn <- neuralnet(case ~ age + parity + induced + spontaneous, data = train,
    hidden = i, stepmax = 10^9, err.fct = "ce", linear.output = FALSE)

  pred <- predict(nn, newdata = train)
  y_hat_train <- round(pred)
  train_err <- mean(train$case != y_hat_train)
  train_err_store <- c(train_err_store, train_err)

  pred <- predict(nn, newdata = test)
  y_hat_test <- round(pred)
  test_err <- mean(test$case != y_hat_test)
  test_err_store <- c(test_err_store, test_err)
}
train_err_store
```

```
## [1] 0.2606061 0.2181818 0.2484848 0.2121212 0.2000000
```

```
test_err_store
```



```
## [1] 0.2650602 0.2168675 0.2409639 0.2771084 0.3012048
```

Neural network model with the outlier

```
set.seed(123)
```

```
train_err_store1 <- c()
test_err_store1 <- c()
head(train_outlier)
```

```
##      education age parity induced case spontaneous stratum pooled.stratum
## 159    12+ yrs 100      2      2    0              0      77           53
## 207     6-11yrs 35      1      0    0              0      42           11
## 179     6-11yrs 29      3      0    0              2      14           29
## 14      6-11yrs 29      3      2    1              0      14           29
## 195     6-11yrs 30      4      1    0              1     30           35
## 170     6-11yrs 35      3      0    0              0       5           32
```

```
for (i in 1:5){

  # fit neural network with "i" neurons
  nn1 <- neuralnet(case ~ age + parity + induced + spontaneous, data = train_outlier,
    hidden = i, stepmax = 10^9, err.fct = "ce", linear.output = FALSE)

  pred1 <- predict(nn1, newdata = train_outlier)
  y_hat_train1 <- round(pred1)
  train_err1 <- mean(train_outlier$case != y_hat_train1)
  train_err_store1 <- c(train_err_store1, train_err1)

  pred1 <- predict(nn1, newdata = test)
  y_hat_test1 <- round(pred1)
  test_err1 <- mean(test$case != y_hat_test1)
  test_err_store1 <- c(test_err_store1, test_err1)
}
train_err_store1
```

```
## [1] 0.2848485 0.2545455 0.2303030 0.2060606 0.1818182
```

```
test_err_store1
```

```
## [1] 0.2409639 0.2168675 0.3132530 0.2771084 0.2650602
```

```
library(knitr)
```

```
data <- data.frame(neurons = 1:5, train_error_without_outlier = train_err_store, train_error_with_outlier = train_err_store1,  
test_error_without_outlier = test_err_store, test_error_with_outlier = test_err_store1)
```

```
kable(data, caption="Model Errors", row.names = FALSE)
```

Model Errors

neurons	train_error_without_outlier	train_error_with_outlier	test_error_without_outlier	test_error_with_outlier
1	0.2606061	0.2848485	0.2650602	0.2409639
2	0.2181818	0.2545455	0.2168675	0.2168675
3	0.2484848	0.2303030	0.2409639	0.3132530
4	0.2121212	0.2060606	0.2771084	0.2771084
5	0.2000000	0.1818182	0.3012048	0.2650602

```
set.seed(123)

# Shrink the outlier towards its original value and retrain the model
shrink_values <- seq(100, 31, length.out = 10)
shrunked_mse <- numeric(length(shrink_values))

for (i in 1:length(shrink_values)) {

  train_outlier$age[1] <- shrink_values[i]

  nn_shrunked <- neuralnet(case ~ age + parity + induced + spontaneous, data = train_outlier, hidden = 2, stepmax = 10^9, er
n.fct = "ce", linear.output = FALSE)

  # Evaluate model performance on the test set
  pred_shrunked <- predict(nn_shrunked, newdata = test)
  y_hat_shrunked <- round(pred_shrunked)
  shrunked_mse[i] <- mean(test$case != y_hat_shrunked)

}

library(knitr)
d <- data.frame(Shrinkage_values = shrink_values, Shrinkage_MSE = shrunked_mse)
kable(d, row.names = FALSE)
```

Shrinkage_values	Shrinkage_MSE
100.00000	0.2650602
92.33333	0.2530120
84.66667	0.2168675
77.00000	0.2891566
69.33333	0.2168675
61.66667	0.3012048

Shrinkage_values	Shrinkage_MSE
54.00000	0.2168675
46.33333	0.2289157
38.66667	0.2289157
31.00000	0.2289157

As shown in the above table the shrink value first stabilizes at 84.66667, mse = 0.2168675

Distance between shrink stabilization and original age value: $85 - 31 = 54$.

Significant changes to network coefficient estimates occur when outlier moves beyond 54 units from original age value(31).

Question 5) Apply a nonlinear SVM to a binary classification data set of your choice. Make up a table of values of Cost and tuning parameter (e.g., gamma). For each cell in the table compute the misclassification rate using cross-validation or a hold out approach. Find the optimal cost and tuning parameter combination from the table. Compare this rate with the results obtained using LDA, logistic regression and a classification tree.

```
library(ISLR2)
library(e1071)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':  
##  
##      select
```

```
## The following object is masked from 'package:ISLR2':  
##  
##      Boston
```

```
data("OJ")  
dim(OJ)
```

```
## [1] 1070  18
```

```
head(OJ)
```

```
## Purchase WeekofPurchase StoreID PriceCH PriceMM DiscCH DiscMM SpecialCH
## 1 CH 237 1 1.75 1.99 0.00 0.0 0
## 2 CH 239 1 1.75 1.99 0.00 0.3 0
## 3 CH 245 1 1.86 2.09 0.17 0.0 0
## 4 MM 227 1 1.69 1.69 0.00 0.0 0
## 5 CH 228 7 1.69 1.69 0.00 0.0 0
## 6 CH 230 7 1.69 1.99 0.00 0.0 0
## SpecialMM LoyalCH SalePriceMM SalePriceCH PriceDiff Store7 PctDiscMM
## 1 0 0.500000 1.99 1.75 0.24 No 0.000000
## 2 1 0.600000 1.69 1.75 -0.06 No 0.150754
## 3 0 0.680000 2.09 1.69 0.40 No 0.000000
## 4 0 0.400000 1.69 1.69 0.00 No 0.000000
## 5 0 0.956535 1.69 1.69 0.00 Yes 0.000000
## 6 1 0.965228 1.99 1.69 0.30 Yes 0.000000
## PctDiscCH ListPriceDiff STORE
## 1 0.000000 0.24 1
## 2 0.000000 0.24 1
## 3 0.091398 0.23 1
## 4 0.000000 0.00 1
## 5 0.000000 0.00 0
## 6 0.000000 0.30 0
```

```
unique(OJ$Purchase)
```

```
## [1] CH MM
## Levels: CH MM
```

```
oj <- OJ
head(oj)
```

```

## Purchase WeekofPurchase StoreID PriceCH PriceMM DiscCH DiscMM SpecialCH
## 1 CH 237 1 1.75 1.99 0.00 0.0 0
## 2 CH 239 1 1.75 1.99 0.00 0.3 0
## 3 CH 245 1 1.86 2.09 0.17 0.0 0
## 4 MM 227 1 1.69 1.69 0.00 0.0 0
## 5 CH 228 7 1.69 1.69 0.00 0.0 0
## 6 CH 230 7 1.69 1.99 0.00 0.0 0
## SpecialMM LoyalCH SalePriceMM SalePriceCH PriceDiff Store7 PctDiscMM
## 1 0 0.500000 1.99 1.75 0.24 No 0.000000
## 2 1 0.600000 1.69 1.75 -0.06 No 0.150754
## 3 0 0.680000 2.09 1.69 0.40 No 0.000000
## 4 0 0.400000 1.69 1.69 0.00 No 0.000000
## 5 0 0.956535 1.69 1.69 0.00 Yes 0.000000
## 6 1 0.965228 1.99 1.69 0.30 Yes 0.000000
## PctDiscCH ListPriceDiff STORE
## 1 0.000000 0.24 1
## 2 0.000000 0.24 1
## 3 0.091398 0.23 1
## 4 0.000000 0.00 1
## 5 0.000000 0.00 0
## 6 0.000000 0.30 0

```

```
set.seed(123)

indis <- sample(1:nrow(oj),size = round(0.65 * nrow(oj)), replace = FALSE)

train_data <- oj[indis, ]
test_data <- oj[-indis, ]

X_train <- train_data[, -1]
Y_train <- train_data[, 1]

X_test <- test_data[, -1]
Y_test <- test_data[, 1]

cost_values <- c(0.01,0.1,1,5,10)
gamma_values <- c(0.01,0.1,0.2,0.5,1)

result_table <- matrix(NA, nrow = length(cost_values), ncol = length(gamma_values))

rownames(result_table) <- as.character(cost_values)
colnames(result_table) <- as.character(gamma_values)

for (cost in cost_values) {
  for (gamma in gamma_values) {

    svm_model <- svm(Purchase ~ ., data = train_data, cost = cost, gamma = gamma)

    pred <- predict(svm_model, newdata = test_data)

    misclassification_rate <- mean(Y_test != pred)

    result_table[as.character(cost), as.character(gamma)] <- misclassification_rate
  }
}

min_misclassification_rate <- min(result_table)
optimal_params <- which(result_table == min_misclassification_rate, arr.ind = TRUE)
optimal_cost <- as.numeric(rownames(result_table)[optimal_params[1]])
optimal_gamma <- as.numeric(colnames(result_table)[optimal_params[2]])
```



```
result_table
```

```
##           0.01      0.1      0.2      0.5      1
## 0.01 0.3903743 0.3903743 0.3903743 0.3903743 0.3903743
## 0.1  0.2139037 0.2112299 0.2566845 0.3021390 0.3556150
## 1    0.1684492 0.1925134 0.2058824 0.2192513 0.2459893
## 5    0.1791444 0.2005348 0.2112299 0.2326203 0.2379679
## 10   0.1764706 0.2005348 0.2165775 0.2352941 0.2433155
```

```
cat("Optimal Cost:", optimal_cost)
```

```
## Optimal Cost: 1
```

```
cat("Optimal Gamma:", optimal_gamma)
```

```
## Optimal Gamma: 0.01
```

```
cat("Min Misclassification Rate:", min_misclassification_rate)
```

```
## Min Misclassification Rate: 0.1684492
```

LDA

```
set.seed(123)
```

```
lda.fit <- lda(Purchase ~., data = train_data)
```

```
## Warning in lda.default(x, grouping, ...): variables are collinear
```

```
lda.fit
```

```
## Call:
## lda(Purchase ~ ., data = train_data)
##
## Prior probabilities of groups:
##      CH      MM
## 0.6106322 0.3893678
##
## Group means:
##   WeekofPurchase StoreID PriceCH PriceMM   DiscCH   DiscMM SpecialCH
## CH      256.4706 4.383529 1.873318 2.099741 0.06811765 0.1000471 0.1835294
## MM      252.2841 3.254613 1.865683 2.063985 0.02671587 0.1674539 0.0996310
##   SpecialMM   LoyalCH SalePriceMM SalePriceCH   PriceDiff Store7Yes   PctDiscMM
## CH 0.1176471 0.7202877   1.999694   1.805200 0.19449412 0.4141176 0.04839651
## MM 0.2287823 0.3119065   1.896531   1.838967 0.05756458 0.1992620 0.08013515
##   PctDiscCH ListPriceDiff   STORE
## CH 0.03593773   0.2264235 1.484706
## MM 0.01407481   0.1983026 1.859779
##
## Coefficients of linear discriminants:
##                               LD1
## WeekofPurchase -0.005596065
## StoreID        -0.028376467
## PriceCH         5.411571785
## PriceMM         5.633107001
## DiscCH          11.039055129
## DiscMM          11.568586519
## SpecialCH       0.292305380
## SpecialMM       0.299973848
## LoyalCH         -3.919138210
## SalePriceMM     -6.678058301
## SalePriceCH     -4.406448590
## PriceDiff       -4.678101073
## Store7Yes       -0.094019783
## PctDiscMM       -45.612963409
## PctDiscCH       -23.423874410
## ListPriceDiff   4.002042780
## STORE          -0.002298879
```

```
test_pred_lda <- predict(lda.fit, newdata = test_data)

table(predicted = test_pred_lda$class , Actual = Y_test)
```

```
##           Actual
## predicted CH  MM
##           CH 200 42
##           MM  28 104
```

```
test_error_lda <- (1/length(test_pred_lda$class))*length(which(Y_test != test_pred_lda$class))
test_error_lda
```

```
## [1] 0.1871658
```

Logistic regression

```
set.seed(123)

glm.fit <- glm(Purchase ~., data = train_data, family = "binomial")
summary(glm.fit)
```

```
##
## Call:
## glm(formula = Purchase ~ ., family = "binomial", data = train_data)
##
## Coefficients: (5 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   5.610558    2.480195   2.262  0.02369 *
## WeekofPurchase -0.008533    0.013257  -0.644  0.51979
## StoreID        -0.015117    0.168593  -0.090  0.92855
## PriceCH         2.519095    2.165983   1.163  0.24482
## PriceMM        -2.628080    1.079881  -2.434  0.01495 *
## DiscCH         13.822294   22.343772   0.619  0.53617
## DiscMM         37.858041   11.244269   3.367  0.00076 ***
## SpecialCH       0.560970    0.422357   1.328  0.18412
## SpecialMM       0.439598    0.341008   1.289  0.19736
## LoyalCH        -6.222700    0.480793  -12.943 < 2e-16 ***
## SalePriceMM           NA           NA      NA      NA
## SalePriceCH           NA           NA      NA      NA
## PriceDiff           NA           NA      NA      NA
## Store7Yes        -0.447960    0.870105  -0.515  0.60667
## PctDiscMM       -74.383869   23.524986  -3.162  0.00157 **
## PctDiscCH       -32.854890   42.006894  -0.782  0.43414
## ListPriceDiff           NA           NA      NA      NA
## STORE            NA           NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 930.50  on 695  degrees of freedom
## Residual deviance: 538.85  on 683  degrees of freedom
## AIC: 564.85
##
## Number of Fisher Scoring iterations: 5
```

```
names(glm.fit)
```

```
## [1] "coefficients"      "residuals"        "fitted.values"
## [4] "effects"           "R"                 "rank"
## [7] "qr"                "family"            "linear.predictors"
## [10] "deviance"          "aic"               "null.deviance"
## [13] "iter"              "weights"           "prior.weights"
## [16] "df.residual"       "df.null"           "y"
## [19] "converged"         "boundary"          "model"
## [22] "call"              "formula"           "terms"
## [25] "data"              "offset"            "control"
## [28] "method"            "contrasts"         "xlevels"
```

```
glm_test <- predict(glm.fit, newdata = test_data, type = "response")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type = if (type ==:
## prediction from rank-deficient fit; attr(*, "non-estim") has doubtful cases
```

```
y_hat_test <- ifelse(glm_test > 0.5, "MM", "CH")
```

```
table(predicted = y_hat_test, Actual = Y_test)
```

```
##           Actual
## predicted CH  MM
##           CH 203 42
##           MM  25 104
```

```
test_err <- mean(Y_test != y_hat_test)
test_err
```

```
## [1] 0.1791444
```

Classification Tree

```
library(rpart)
library(rpart.plot)


set.seed(123)

model.control <- rpart.control(minsplit = 3, xval = 10, cp = 0)
fit.X <- rpart(Purchase ~., data = train_data, method = "class", control = model.control)
names(fit.X)
```

```
## [1] "frame"          "where"          "call"
## [4] "terms"          "cptable"        "method"
## [7] "parms"          "control"        "functions"
## [10] "numresp"        "splits"         "csplit"
## [13] "variable.importance" "y"              "ordered"
```


```
plot(fit.X$cptable[,4], main = "Xval error for model selection", ylab = "xerr",pch=19,type
='o')
```

Xval error for model selection



```
plot(fit.X$cptable[,1], main = "Cp for model selection", ylab = "Cp",pch=19,type='o')
```

Cp for model selection




```
min_cp <- which.min(fit.X$cptable[,4])  
min_cp
```

```
## 4  
## 4
```

```
pruned_fit_X <- prune(fit.X, cp = fit.X$cptable[min_cp, 1])  
  
plot(fit.X, branch = .3, compress = T, main = "Full Tree")  
text(fit.X, cex = .7)
```



Full Tree



```
rpart.plot(fit.X,digits = 3, extra = 1,main = "Full Tree")
```


```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```

Full Tree



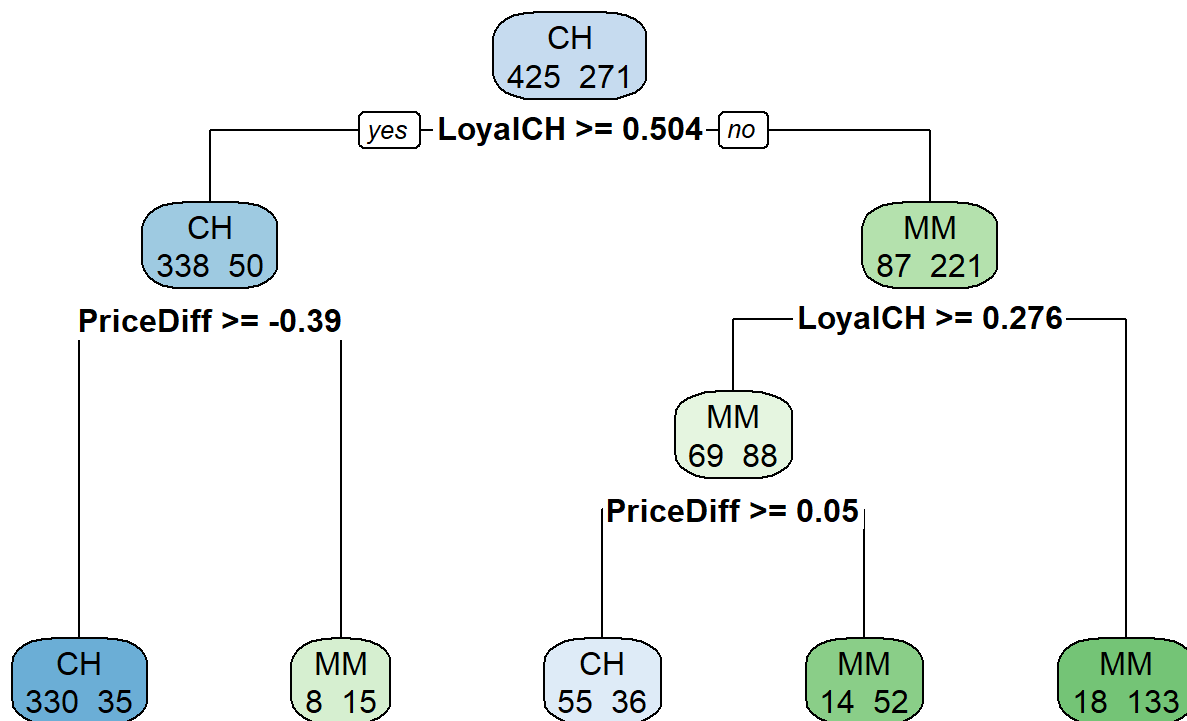
```
plot(pruned_fit_X, branch = .3, compress = T, main = "Pruned Tree")
text(pruned_fit_X, cex = .7)
```

Pruned Tree



```
rpart.plot(pruned_fit_X,digits = 3, extra = 1,main = "Pruned Tree")
```

Pruned Tree



```
pred_test <- predict(pruned_fit_X, newdata = test_data, type = 'class')
table(predicted = pred_test, Actual = Y_test)
```

```
##      Actual
## predicted CH  MM
##      CH 207  50
##      MM  21  96
```

```
err <- sum(pred_test != Y_test) / nrow(test_data)
cat("Misclassification rate for a Classification Tree :", err * 100, "\n")
```

```
## Misclassification rate for a Classification Tree : 18.98396
```

Comparison between the models

```
library(dplyr)
```

```
data <- data.frame(Methods = c("Non linear SVM", "LDA", "Logistic Regression", "Classification Tree"), misclassification_rate = c(min_misclassification_rate * 100, test_error_lda * 100, test_err * 100, err * 100))
```

```
arrange(data, misclassification_rate)
```

```
##           Methods misclassification_rate
## 1   Non linear SVM           16.84492
## 2 Logistic Regression           17.91444
## 3             LDA           18.71658
## 4 Classification Tree           18.98396
```

On the OJ data set Non-linear SVM is performing well with low test mse = 16.844 when compared with the other models.

Next logistic regression performing better with test mse = 17.914

LDA model is performing with test mse = 18.716

And Classification Tree performance is poor compared to the other models on this OJ data set.